

Integration of AWS Redshift with CADP

AWS Redshift [Overview]

This document describes how to configure and integrate CipherTrust Manager with AWS Redshift.

Amazon Redshift is a data warehouse product which forms part of the larger cloud-computing platform Amazon Web Services. It is built on top of technology from the massive parallel processing (MPP) data warehouse company ParAccel to handle large scale data sets and database migrations. Redshift differs from Amazon's other hosted database offering, Amazon RDS, in its ability to handle analytic workloads on big data data sets stored by a column-oriented DBMS principle. Redshift allows up to 16 petabytes of data on a cluster compared to Amazon RDS Aurora's maximum size of 128 terabytes. Thales provides a couple of different methods to protect sensitive data in AWS Redshift.

Bring Your Own Encryption (BYOE)

- **Data Ingest** – with Thales Batch Data Transformation (BDT)
- **Data Access** –user defined functions for column level encryption and decryption using Thales CADP and tokenization using Thales CT-VL.

Bring/Hold Your Own Key (BYOK) (HYOK)

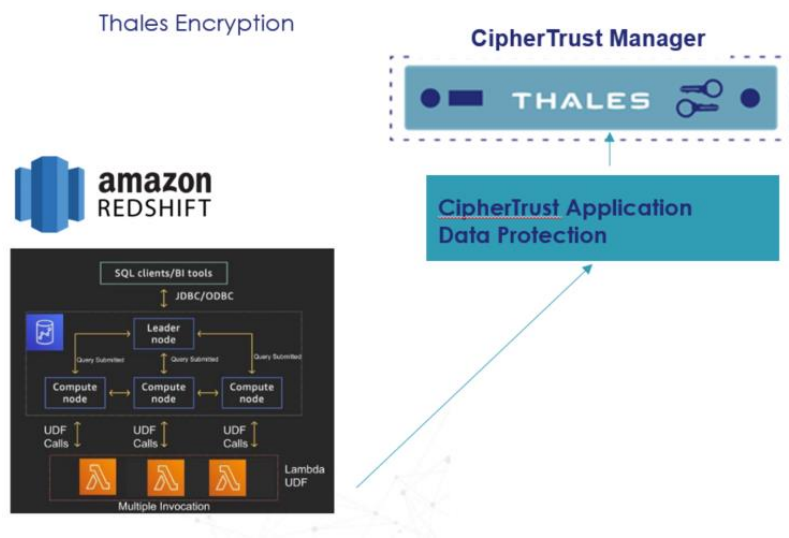
- **AWS Redshift Customer Managed Keys-** with Thales CM CCKM BYOK and HYOK.

The above methods are NOT mutually exclusive. All methods can be used to build a strong defense in depth strategy to protect sensitive data in the cloud. The focus of this integration will be on Data Access protecting sensitive data in AWS Redshift columns by using CADP to create User Defined Functions (UDF) for encryption and decryption of sensitive data.

Architecture

The examples provided in this document use a capability AWS Redshift called “External Function”. A Redshift external function lets you incorporate functionality with software outside of Redshift by providing a direct integration with AWS Lambda Functions. A Redshift external function allows you to implement your function in other languages other than SQL. Listed

below is a diagram of how this integration works.



Supported Product Versions

- **CipherTrust Manager** CipherTrust Manager 2.11 and higher. **Note:** If usersets are enabled then CipherTrust Manager should be at 2.14 or above.
- **CADP** CADP Java 8.13 and higher. **Note:** CADP 8.15.0.001 was tested. If a higher version of CADP is used you will encounter a class not found error with AWS Lambda Functions.
- **AWS Redshift**

This integration is validated using 2nd generation Google Cloud Functions and Java 11 along with CM 2.14 and CADP 8.15.0.001

Prerequisites

Steps performed for this integration were provided by this AWS link:

<https://docs.aws.amazon.com/redshift/latest/dg/udf-creating-a-lambda-sql-udf.html>

https://docs.aws.amazon.com/redshift/latest/dg/r_CREATE_EXTERNAL_FUNCTION.html

- Ensure that CADP for Java is installed and configured. Refer to <https://thalesdocs.com/ctp/con/cadp/cadp-java/latest/admin/cadp-for-java-quick-start/cadp-for-java-installer/index.html>
- Ensure that the CipherTrust Manager is installed and configured. Refer to the [CipherTrust Manager documentation](#) for details.
- AWS Lambda functions communicates with the CipherTrust Manager using the Network Attached Encryption (NAE) Interface. Ensure that the NAE interface is configured. Refer to the [CipherTrust Manager documentation](#) for details.

Steps for Integration

- [Installing and Configuring Thales CADP for Java]
- [Download code from Thales github and compile]
- [Publish jar/zip file to AWS Lambda Function (endpoint)]
- [Create AWS Redshift External Function]
- [Integration with Thales CipherTrust Manager]
- [Environment Variables]

Installing and Configuring CADP for Java

To install and configure **CADP for Java**, refer to [Quick Start](#).

As noted in the link above Thales provides CADP in the Maven repository. This can be found at: https://central.sonatype.com/artifact/io.github.thalescpl-io.cadp/CADP_for_JAVA

Note: As indicated earlier please use the 8.15.0.001 version with Redshift. Eclipse development tool was used for these examples. Here is the version used for testing along with the Maven plugin for Eclipse.

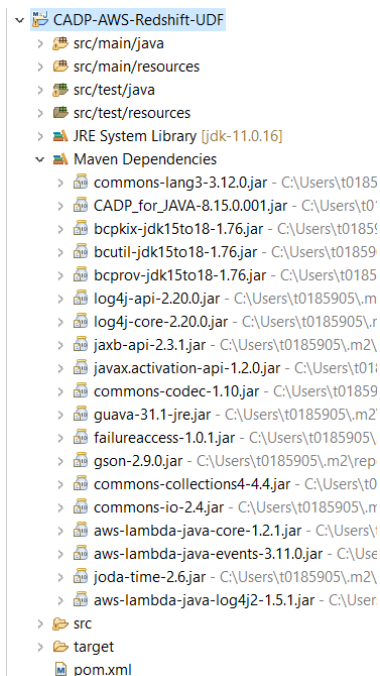
```
eclipse.buildId=4.15.0.I20200305-0155
```

```
m2e - Maven Integration for Eclipse
```

Example:

```
<dependency>
  <groupId>io.github.thalescpl-io.cadp</groupId>
  <artifactId>CADP_for_JAVA</artifactId>
  <version>8.15.0.001</version>
</dependency>
```

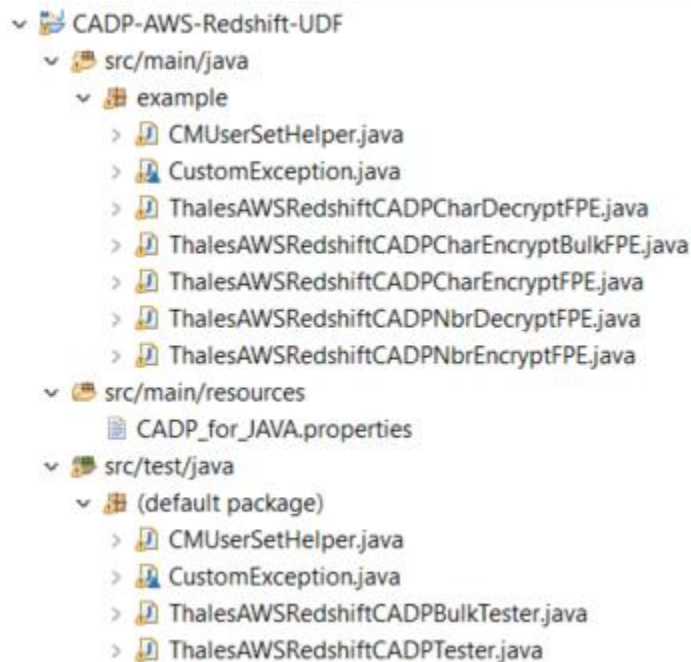
Here is a screenshot in eclipse of the pom file used for these examples:



Download code from github and compile.

git clone https://github.com/ThalesGroup/CipherTrust_Application_Protection.git

The database directory has all the code for AWS Redshift. You should see the classes in your project.



CADP supports a bulk API which allows for CADP to batch requests before calling encrypt or decrypt. A separate class file for each datatype is available for testing with this API. There is a limit of 10,000 items to encrypt but it also depends on the size of the items as well so review the CADP documentation to make sure those thresholds are not exceeded. Typically, the point of diminishing returns is around 200.

Assuming you have your CM already configured the `ThalesAWSRedshiftCADPTester` can be used to test basic connection to CM to make sure your CM environment is configured correctly. You will need to modify the keyname to make sure it exists in CM and there are environment variables that must be set as well. If you plan on using UserSets the user must also be in the Application Data Protection Client Group. See section on Environment Variables for more details. If you have already installed CM then you will need to update the `CADP_for_JAVA.properties` file with all the necessary settings such as IP/NAE Port, etc. The file is located under the resource's directory in the java project for eclipse.

Generate the jar file to upload to the CSP.

To compile and generate the target jar file to be uploaded to AWS Lambda select the project and choose "Run As" "maven install" to generate the target.

```
[INFO] Replacing C:\Users\t0185905\workspace\CADP-AWS-Redshift-UDF\target\CADP-AWS-Redshift-UDF-0.0.2-SNAPSHOT.jar with C:\Users\t0185905\workspace\CADP-AWS-Redshift-UDF\target\CADP-AWS-Redshift-UDF-0.0.2-SNAPSHOT-shaded.jar
[INFO] --- maven-install-plugin:3.0.1:install (default-install) @ CADP-AWS-Redshift-UDF ---
[INFO] Installing C:\Users\t0185905\workspace\CADP-AWS-Redshift-UDF\pom.xml to C:\Users\t0185905\.m2\repository\Thales\CADP-AWS-Redshift-UDF\0.0.2-SNAPSHOT\CADP-AWS-Redshift-UDF-0.0.2-SNAPSHOT.pom
[INFO] Installing C:\Users\t0185905\workspace\CADP-AWS-Redshift-UDF\target\CADP-AWS-Redshift-UDF-0.0.2-SNAPSHOT.jar to C:\Users\t0185905\.m2\repository\Thales\CADP-AWS-Redshift-UDF\0.0.2-SNAPSHOT\CADP-AWS-Redshift-UDF-0.0.2-SNAPSHOT.jar
[INFO] -----
[INFO] BUILD SUCCESS
[INFO] -----
[INFO] Total time: 4.191 s
[INFO] Finished at: 2024-05-15T15:59:06-04:00
[INFO] -----The
```

code provided has uses Googles userDefinedContext capability. This code accepts a mode and a datatype as a keyvalue pair.

Publish jar/zip file to AWS Lambda Function. (endpoint)

Once you have generated the jar file to upload you can then create the CSP function.

AWS Lambda Function

Create function [Info](#)

Choose one of the following options to create your function.

☒ **Author from scratch**
 Start with a simple Hello World example.

☐ **Use a blueprint**
 Build a Lambda application from sample code and configuration presets for common use cases.

☐ **Container image**
 Select a container image to deploy for your function.

Basic information

Function name
 Enter a name that describes the purpose of your function.

 Use only letters, numbers, hyphens, or underscores with no spaces.

Runtime [Info](#)
 Choose the language to use to write your function. Note that the console code editor supports only Node.js, Python, and Ruby.

Architecture [Info](#)
 Choose the instruction set architecture you want for your function code.
☒ x86_64
☐ arm64

Permissions [Info](#)
 By default, Lambda will create an execution role with permissions to upload logs to Amazon CloudWatch Logs. You can customize this default role later when adding triggers.

[Change default execution role](#)

[Advanced settings](#)

[Cancel](#) [Create function](#)

Click Create function.

Set environment variables appropriate values (See section on Environment Variables for details) and then select Next. Upload the jar or zip file on the next screen. Be sure to change the entry point to reflect your class name. ***The example below is*** `ThalesAWSRedshiftCADPNbrDecryptFPE` ***This should match your code path.*** Be sure to have a role that allows for `AWSLambdaBasicExecutionRole`. 512MB of RAM is plenty. You can monitor your tests and possibly lower it.

Once you have created the functions above and if you have already configured and setup CM with the key and all the environment variables you can test the function with the test tab. You will need to provide the appropriate json to test. Here is an example:

```
{
  "request_id": "23FF1F97-F28A-44AA-AB67-266ED976BF40",
  "cluster": "arn:aws:redshift:xxxx",
  "user": "adminuser",
  "database": "db1",
  "external_function": "public.foo",
  "query_id": 5678234,
  "num_records": 3,
  "arguments": [
    [
      "The. king. and I me"
    ],
    [

```

```
    "The mroe data "
  ],
  [
    "test dada"
  ]
]
```

Create AWS Redshift External Function.

Here are some links that provide details.

<https://docs.aws.amazon.com/redshift/latest/dg/udf-creating-a-lambda-sql-udf.html>

https://docs.aws.amazon.com/redshift/latest/dg/r_CREATE_EXTERNAL_FUNCTION.html

<https://aws-dojo.com/excercises/excercise31/>

As noted above the steps are:

1. Create the AWS Lambda Function. (should already be done from above)
2. Create external function object in AWS Redshift

Examples:

Redshift Function Definitions.

The functions are created using the following format:

```
CREATE OR REPLACE EXTERNAL FUNCTION thales_token_cts_char(sensitivedata
varchar)
RETURNS varchar
STABLE
LAMBDA 'thales_aws_lambda_redshift_cts_token_char'
IAM_ROLE 'arn:aws:youriam/hiftpoweruseraccess';
```

Integration with CipherTrust Manager.

Assuming the NAE interface has already been setup based on the prerequisites the only other setup required is to have a user and key created in CM. This user should have the ability to export the key. This user will also be needed to be provided as an environment variable for the function. The examples provided have the key as a hardcoded value, but this can be easily altered to be provided as an environment variable, obtained from a secrets manager.

As noted above there is a test class that can be used to test connectivity with CM without having to publish the Function. If you have not updated the CADP_for_JAVA.properties file with the CM settings such as IP/NAE Port, etc. then do so now. The file is located under the resource's directory in the eclipse project. These properties can also be overwritten with CADP code as well if your desire is to pass them in as environment variables or headers of the JSON request. Here is an example:

```
System.setProperty("com.ingrian.security.nae.NAE_IP.1", "10.20.1.9");
```

When all the above steps are performed you should see your UDF's in AWS Redshift under Routines in the UI. Here is a sample query using one of the UDF's.

Sample Results:

The screenshot displays the AWS Redshift Query Editor v2 interface. On the left, the 'warner-redshift-cluster-1' resource is expanded, showing the 'public' schema with a list of tables including 'category', 'date', 'event', 'listing', 'sales', 'users', and 'venue'. The 'Functions' section is also expanded, showing three user-defined functions: 'thales_cadp_decrypt_char(varchar)', 'thales_token_cadp_char(varchar)', and 'thales_token_ots_char(varchar)'. The main query editor shows a SQL query: `select thales_cadp_decrypt_char(eventnameenc) as decvalue, eventname, eventnameenc from event_enc_test`. The results are displayed in a table with 40 rows. The first row shows 'The King and I' for both eventname and eventnameenc, with a decvalue of '8dh Oc0l qMK 9'. Subsequent rows show other event names and their corresponding encrypted values and decrypted values.

decvalue	eventname	eventnameenc
The King and I	The King and I	8dh Oc0l qMK 9
The Bacchae	The Bacchae	Fot RklqRGf
Herbie Hancock	Herbie Hancock	fg65e2 zZtlqV
Gretchen Wilson	Gretchen Wilson	Kvy3Dv8O IpoHt0
Dierks Bentley	Dierks Bentley	bkrWGX aYSFcTE
Young Frankenstein	Young Frankenstein	QlnAb K3gytKLnPNdf
The Caucasian Chalk Circle	The Caucasian Chalk Circle	86m ZFBT2Dmdl M2KHF ...
King Crimson	King Crimson	acjD T5wyYMM
My Morning Jacket	My Morning Jacket	jq GEholrP 7XyExn
George Thorogood and th...	George Thorogood and th...	k1xeAm LAZfmIrbB okH x...
Alkaline Trio	Alkaline Trio	A04QalfH5 snLg
Counting Crows and Mar...	Counting Crows and Mar...	9guXWet.0 Lg1tQ Tau Re...
Jersey Boys	Jersey Boys	tVUDrc h65w
The Caucasian Chalk Circle	The Caucasian Chalk Circle	86m ZFBT2Dmdl M2KHF ...
Boyz Scaggs	Boyz Scaggs	D5r cmuq7T
Missy Higgins	Missy Higgins	3ooJg ebLqMvH
LeAnn Rimes	LeAnn Rimes	Uedhn jDc1
Fab Faux	Fab Faux	NULL

Environment Variables

Listed below are the environment variables currently required for the Cloud Function.

```
String userName = System.getenv("CMUSER");
// If you plan on using UserSets the cmuser must also be in the Application Data
// Protection Client Group.
String password = System.getenv("CMPWD");
// returnciphertextforuserwithnokeyaccess = is na environment variable to express
```



```
// how data should be returned when the user above does not have access to the key
//and if doing a lookup in the userSet and the user does not exist. If
//returnCiphertextForUserWithNoKeyAccess = null
// then an error will be returned to the query, else the results set will provide
ciphertext.
// validValues are 1 or null
// 1 will return cipher text
// null will return error.
String returnCiphertextForUserWithNoKeyAccess =
System.getenv("returnCiphertextForUserWithNoKeyAccess");
// userSetLookup = should a userSet lookup be done on the user from Big Query? 1
// = true 0 = false.
String userSetLookup = System.getenv("userSetLookup");
// userSetID = should be the userSetId in CM to query.
String userSetID = System.getenv("userSetIdInCM");
// userSetLookupIP = this is the IP address to query the userSet. Currently it
// is the userSet in CM but could be a memcache or other in memory db.
String userSetLookupIP = System.getenv("userSetLookupIP");
```

Advanced Topics.

Google also provides the ability to setup a secure perimeter with VPC Service Controls. See link below for more information.

https://cloud.google.com/Redshift/docs/remote-functions#using_vpc_service_controls

Application Data Protection UserSets

Application Data Protection UserSets are currently used for DPG to control how the data will be revealed to users. These UserSets can also be independent of any Access Policy. Most cloud databases have some way to capture who is running the query and this information can be passed to CM to be verified in a UserSet to ensure the person running the query has been granted proper access. In github there is a sample class file called CMUserSetHelper that can be used to load a userSet with values from an external identity provider such as LDAP. The name of this method is `addAUserToUserSetFromFile`. Once users have been loaded into this userSet the userSetId must be captured and used as an environment variable to the Function. The function has a number of environment variables that must be provided for the function to work. Please review the section on Environment Variables for more details.

Options for keyname

Currently the sample code uses a hard coded key name. Other options to be considered are:

- 1.)keyname as an environment variable.
- 2.)keyname as a AWS secret.

3.)keyname passed in an attribute to the function. Example with a database view.

create view employee as select first_name, last_name,

thales_cadp_aws_decrypt_char('testfaas', email) as email from emp_basic

Options for handling null values.

Since it is not possible to encrypt a column that contains null values or any column that has 1 byte it is necessary to skip those to avoid getting an error when running the query. There are a couple of ways to handle this use case.

Option 1. Modify the queries.

Many times, simply adding a where clause to exclude values that have nulls or less than 2 bytes can avoid query errors. For example: select * from FROM

mw_demo_dataset_US.plaintext50cadpemailprotected_nulltest where email is not null and length(email) > 1; For those scenarios where that is not suffice some other examples are listed below. Here is an example of a select statement that can be modified to handle null values:

Here is an example of a select statement that can be modified to handle null values:

```
SELECT
  name,
CASE
  WHEN email IS NULL THEN 'null'
  WHEN length(email) < 2 then email
  WHEN email = 'null' then email
  ELSE `cpl-sales-1-app-us-01.mw_demo_dataset_US.thales_cadp_encrypt_char`(email)
END AS email
FROM
  mw_demo_dataset_US.plaintext50cadpemailprotected_nulltest;
```

The above use case was for situations where the column contained both null and the word 'null'.

Here is an example of a select statement that can be modified to handle null values in the where clause.

```
SELECT name, email, email_enc
FROM
mw_demo_dataset_US.plaintext50cadpemailprotected_nulltest
where CASE
  WHEN email_enc IS NULL THEN 'null'
  WHEN length(email_enc) < 2 then email_enc
  WHEN email_enc = 'null' then email_enc
  ELSE `cpl-sales-1-app-us-01.mw_demo_dataset_US.thales_cadp_decrypt_char`(email_enc)
END like "%gmail%"
```

Row	name	email	email_enc
1	Dr. Lemmie Zboncak	ikris@gmail.com	1IOEk@RPDqC.GHd
2	Zillah Leuschke	scronin@gmail.com	53mWY7Y@4bzb6.2D4
3	Troy Gaylord	devon49@gmail.com	EsTMziF@ISZg2.yN6
4	Dr. Spurgeon Wintheiser	hilah17@gmail.com	PIM4vAd@qAIV9.oJC
5	Tatyana Bernhard	denny56@gmail.com	yVQ7ITA@3idYW.GqV
6	Renata Hilpert	tdickens@gmail.com	plzg3zLb@oetcj.Opi
7	Deliah Douglas	sconnelly@gmail.com	jpXAUZWpX@koaH2.yS8
8	Amare Feeney	batz.geary@gmail.com	G04W.CRNlb@riDnQ.DHT
9	Dr. Cristy Schinner	schulist.garfield@gmail.com	Hll90wCU.LPtU3p6o@F9Nb7.G...
10	Vena Douglas	huston.christiansen@gmail.com	7OHpfz.PXI0PoJJJaWlQ@AvTP...

Option 2. Modify the Cloud Function to skip encrypting these values.

Include an if statement checking for null values the word 'null' or any value less than 2.

```

if (redshiftrow != null && redshiftrow.size() > 0) {
    JsonElement element = redshiftrow.get(0);
    if (element != null && !element.isJsonNull()) {
        sensitive = element.getAsString();
        if (sensitive.isEmpty() || sensitive.length() < 2) {
            if (sensitive.isEmpty()) {
                JsonElement elementforNull = new JsonPrimitive("null");
                sensitive = elementforNull.getAsJsonPrimitive().toString();
            } else {
                sensitive = element.getAsJsonPrimitive().toString();
            }
            encdata = sensitive;
        } else {
            if (sensitive.equalsIgnoreCase("null")) {
                sensitive = element.getAsJsonPrimitive().toString();
                encdata = sensitive;
            } else {
                sensitive = element.getAsJsonPrimitive().toString();
                byte[] outbuf = encryptCipher.doFinal(sensitive.getBytes());
            }
        }
    } else {
        JsonElement elementforNull = new JsonPrimitive("null");
        sensitive = elementforNull.getAsJsonPrimitive().toString();
        encdata = sensitive;
    }
} else {
    JsonElement elementforNull = new JsonPrimitive("null");
    sensitive = elementforNull.getAsJsonPrimitive().toString();
    encdata = sensitive;
}

```

Note: When using this Logic it is important to know that for any column that is null will return 'null'. This should be fine for use cases where the query is not updating the column. For use cases where the source system is expecting null vs the word 'null' additional testing should be conducted on the systems that rely on this data type as being null vs the word 'null'.